


RAG Architecture: Version 1 vs. Version 2 Multimodal PDF Pipeline

High-level ingestion, chunking, retrieval, and answer-generation architecture

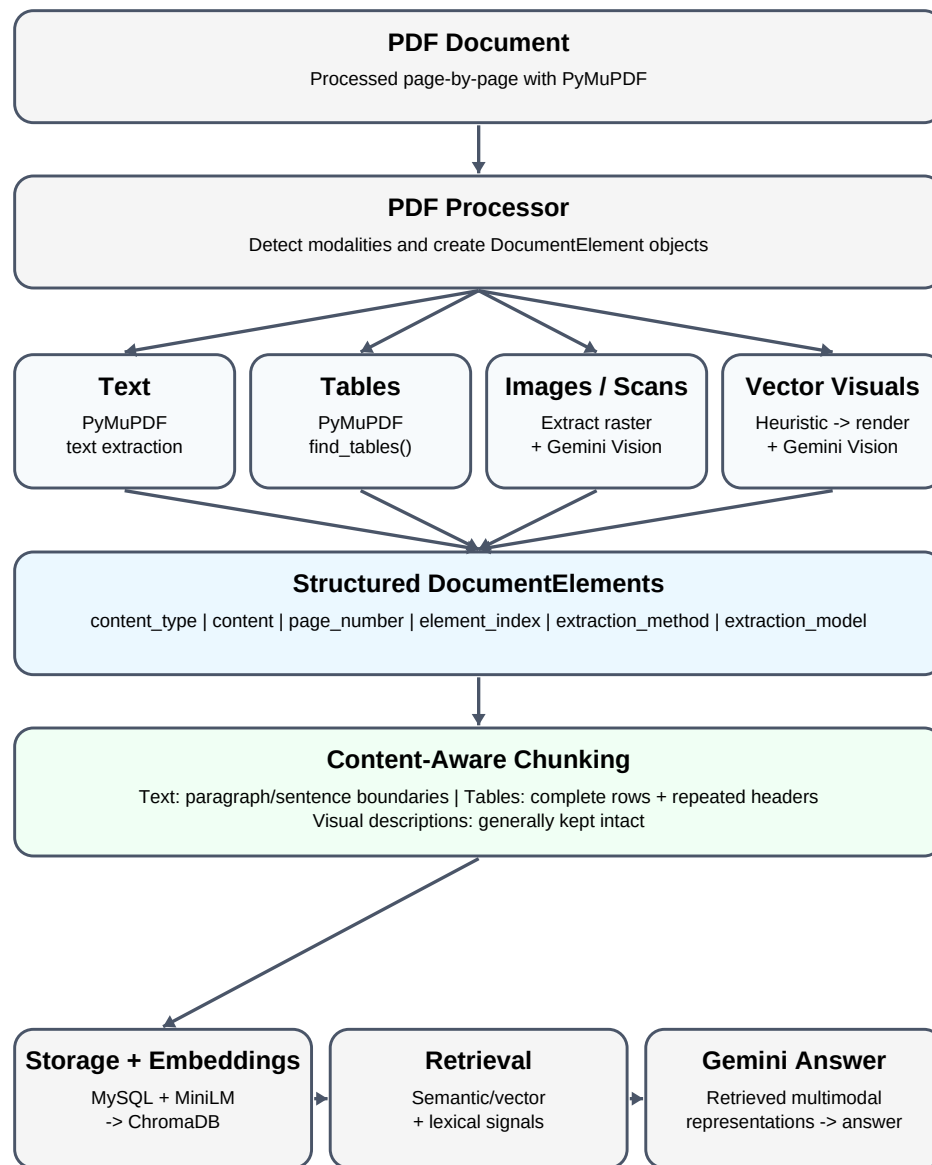
Version 1 - Text-Centric Pipeline



V1 limitation:

Visual PDF information is not reliably represented in retrieval.

Version 2 - Multimodal PDF Pipeline



Current V2 scope: multimodal/content-aware processing applies to PDFs; TXT/DOCX retain V1 chunking.